A Dual-Path Neuro-Symbolic Verification Framework for Reliable Reasoning with Frozen Large Language Models

Anonymous Author(s)
Submitted for double-blind review

***Abstract*—Large language models (LLMs) remain unreliable on tasks that require strict logical consistency, and fine-tuning is often impractical when the base model must stay frozen for cost, licensing, or auditability reasons. We study whether an independent neural-to-symbolic verification branch, operating on the same raw natural-language input as a frozen LLM, can improve that LLM's logical reliability without modifying its parameters. We build a dual-path architecture in which a frozen Qwen3-8B produces an initial answer while, in parallel, the official NL2Logic framework – using Gemma-2-9B-Instruct as its recursive semantic-parsing model – translates the same raw text into a first-order-logic abstract syntax tree that is deterministically compiled into Z3 and solved. A deterministic verifier then retains, corrects, or abstains from the neural prediction based on agreement with the symbolic verdict. On a frozen 20-item raw-text ProofWriter pilot, the dual-path system reaches 70.0% accuracy (macro-F1 0.676) versus 40.0% (macro-F1 0.300) for Qwen3-8B alone, with all six observed neural/symbolic disagreements resolved correctly by the symbolic branch and none harmed. On a parallel five-item FOLIO gate, however, the same bridge produces syntactically valid, Z3-executable theories in 96.6% of sentences but only 1 of 5 correct downstream verdicts, because cross-sentence entity and relation linking breaks down on FOLIO's more natural language. We report this FOLIO result as a generalization limitation rather than a negative result about NL2Logic in general, and conclude that neural-to-symbolic semantic translation – not the symbolic solver – remains the principal bottleneck for extending this verification strategy beyond controlled, closed-vocabulary language.**

***Index Terms*—neuro-symbolic reasoning, large language models, first-order logic, Z3, logical verification, frozen models**

I. Introduction

Large language models (LLMs) are increasingly used for tasks that require multi-step logical reasoning, yet their answers are produced by next-token prediction rather than by any process that guarantees logical consistency. Even strong instruction-tuned models can accept an unsatisfiable set of premises, ignore a negation, or silently misapply a quantifier, and because the failure is embedded in fluent text it is easy to miss. Fine-tuning a model to be more reliable is not always an option: the model may be served behind an API, shared across products, subject to a fixed evaluation protocol, or simply too expensive to retrain every time a new failure mode is found. This motivates verification strategies that treat the LLM as *frozen* and instead check its output against an independent source of ground truth.

Symbolic reasoning is an attractive independent check because a solver such as Z3 is deterministic and its correctness does not depend on training data. The obstacle is translation: converting raw natural language into a formal representation that a solver can consume is itself an open, LLM-dependent problem, and a translation bridge that produces syntactically valid but semantically wrong logic can make a verification system worse, not better. Recent work such as NL2Logic [1] addresses this by decoupling semantic parsing (an LLM builds a first-order-logic abstract syntax tree, one clause at a time) from code generation (a deterministic two-pass compiler emits solver-ready code from that tree), reporting large gains in both syntactic executability and downstream semantic accuracy over unconstrained LLM translation.

This paper asks a narrower, system-level question: *given* a bridge like NL2Logic, does wiring it into an independent verification branch alongside a frozen LLM actually improve that LLM's reliability on raw-text logical reasoning, and where does the resulting system break? We build a dual-path architecture – Neuro-Symbolic Cross-Encoding (NSCE) – in which the same raw premises and hypothesis are given to (i) a frozen Qwen3-8B [2] and (ii) an independent bridge consisting of NL2Logic with Gemma-2-9B-Instruct [3] as its parsing model, followed by Z3 [4]. A deterministic verifier compares the two verdicts and decides to RETAIN the neural answer, CORRECT it with the symbolic verdict, or ABSTAIN when the bridge itself fails. We evaluate this system on ProofWriter [5] and FOLIO [6], chosen because they represent, respectively, controlled synthetic language and more natural crowd-sourced language, and because both are used in NL2Logic's own evaluation.

Our contributions are:

- An independent dual-path verification architecture in which the same raw natural-language evidence is processed by a frozen neural LLM and an independent neural-to-symbolic-plus-solver branch, with an explicit RETAIN/CORRECT/ABSTAIN decision rather than a hard swap to the symbolic answer.
- An empirical study, run under a predefined go/no-go gating protocol, showing a +30 percentage point accuracy improvement (40.0% $\rightarrow$ 70.0%) over a frozen Qwen3-8B on a raw-text ProofWriter pilot, with zero harmful corrections observed.
- A parallel study on FOLIO showing that the identical bridge, on the same hardware and models, achieves near-perfect syntactic executability (96.6%) but only 1/5 cor-

rect downstream verdicts, isolating cross-sentence semantic translation – not solver reliability – as the dominant failure mode when moving from controlled to more natural language.
- A quantitative breakdown of verifier behavior (retain/correct/abstain counts, helpful vs. harmful corrections, bridge-failure rate, and latency) intended to make the limits of this style of system, and not only its headline number, reproducible and legible.

We emphasize what this paper does *not* claim. We do not claim to have invented neuro-symbolic reasoning, NL-to-FOL translation, AST-guided parsing, or the use of Z3 for entailment checking; all of these are established prior work, and NL2Logic in particular is used here strictly as an external, unmodified dependency. Our contribution is the verification-branch architecture built around it, and the empirical characterization of where it succeeds and where it does not.

II. Related Work

Neuro-symbolic reasoning with LLMs. Combining LLMs with symbolic components to obtain formally checkable reasoning is an active area [7]. Logic-LM [8] translates natural-language premises and a hypothesis into first-order logic (or other formal languages) using few-shot prompting and delegates inference to an external solver, optionally applying a self-refinement loop when the generated program fails to compile. LINC [9] similarly combines LLM-generated first-order logic with an automated theorem prover, majority-voting over multiple translation samples to improve robustness. Both systems treat translation as a largely single-step, free-form generation task, which – as later work documents – constrains achievable syntactic and semantic accuracy.

Constrained and structured NL-to-FOL translation. Grammar-Constrained Decoding (GCD) [10] enforces token-level syntactic validity by restricting an LLM's decoding to a context-free grammar for the target formal language, improving executability but not guaranteeing that a grammatically valid formula is semantically faithful to the source sentence. CODE4LOGIC [11] instead prompts a code-generation model to emit Python code encoding a grammar tree, executing that code to obtain the final formula. NL2Logic [1], which we build directly on, decouples semantic parsing from code generation via an explicit first-order-logic abstract syntax tree (FOLAST): a recursive, clause-level LLM parser constructs the tree, and a deterministic two-pass compiler – independent of the LLM – emits solver-specific code (e.g., Z3) from it. The original NL2Logic evaluation reports near-perfect syntactic correctness and an average 30% semantic-accuracy improvement over GCD across thirteen models on FOLIO, ProofWriter, and LogicNLI [12], and a 31% downstream accuracy improvement when substituted into Logic-LM's translation module.

Reasoning-time scaffolding without formal verification. Chain-of-thought prompting [13] improves LLM reasoning accuracy by eliciting intermediate steps but provides no independent, checkable proof that those steps are valid – a limitation this line of work shares with any approach that treats the LLM itself as the sole source of ground truth.

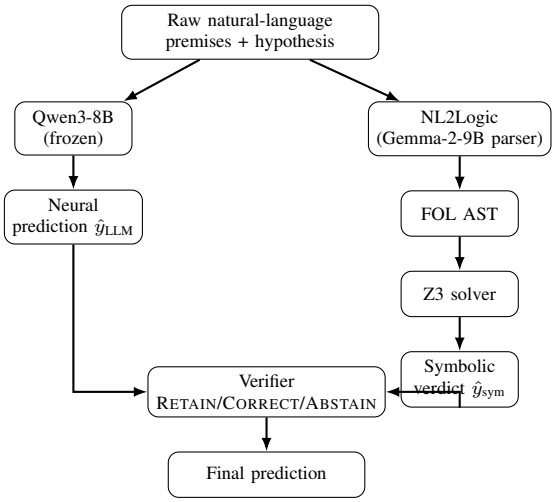


Fig. 1. The dual-path NSCE architecture. The same raw text feeds a frozen neural branch (Qwen3-8B) and an independent neural-to-symbolic branch (NL2Logic with Gemma-2-9B, compiled to Z3); a deterministic verifier reconciles the two verdicts.

Positioning of this work. We do not propose a new NL-to-FOL translator, a new solver, or a new benchmark. We instead study NL2Logic *as an independent verification bridge for a separate, frozen LLM*, rather than as the primary predictor it was designed and evaluated to be in its own publication. This shifts the relevant questions from "how good is the translator in isolation" (already answered by [1]) to "does wiring an existing translator into a verification loop measurably help a frozen model, and does that benefit survive contact with more natural language."

III. Proposed NSCE Architecture

Fig. 1 shows the overall system. The same raw natural-language premises and hypothesis are given, unmodified, to two independent branches.

A. Neural Reasoning Branch

A frozen Qwen3-8B [2] instance receives the raw premises and hypothesis through a fixed prompt template and produces one of TRUE, FALSE, or UNKNOWN via constrained greedy decoding restricted to those three label tokens. The model is never fine-tuned, and no gold label or gold logical form is ever included in its prompt.

B. Neural-to-Symbolic Bridge

The same raw premises and hypothesis, sentence by sentence, are passed to the official NL2Logic framework [1], using Gemma-2-9B-Instruct [3] as the semantic-parsing LLM inside NL2Logic's own recursive pipeline. For each sentence, NL2Logic's `ParserSelector` classifies it as atomic, quantified, or logical (connective-bearing), then recursively dispatches to a specialized sub-parser that extracts the outermost logical construct and re-feeds any remaining sub-sentence back

Algorithm 1 NSCE verification and feedback

Require: raw premises p, raw hypothesis h
$\hat{y}_{\text{LLM}} \leftarrow$ Qwen3-8B(p, h) {Branch A, frozen}
ast $\leftarrow$ NL2Logic$_{\text{Gemma-2-9B}}(p, h)$ {Branch B}
if translation or Z3 compilation of ast fails **then**
  return $\hat{y}_{\text{LLM}}$, status $\leftarrow$ ABSTAIN
end if
$\hat{y}_{\text{sym}} \leftarrow$ SOLVEZ3(ast) {Eq. (1)}
if $\hat{y}_{\text{LLM}} = \hat{y}_{\text{sym}}$ **then**
  return $\hat{y}_{\text{LLM}}$, status $\leftarrow$ RETAIN
else
  return $\hat{y}_{\text{sym}}$, status $\leftarrow$ CORRECT
end if

into the parser, until each sentence is reduced to a first-order-logic abstract syntax tree (FOLAST). This process is entirely NL2Logic's own design; we neither modify nor reimplement it.

C. Symbolic Reasoning Layer

Each sentence's FOLAST is compiled, via NL2Logic's deterministic two-pass generator, into Z3-executable declarations and expressions. We accumulate the per-sentence declarations and assertions for all premises of a given item into one shared Z3 solver context, and check entailment of the hypothesis H against the conjoined premises P using the standard three-way test:

$$\text{verdict}(P, H) = \begin{cases} \text{TRUE}, & \text{UNSAT}(P \wedge \neg H) \\ \text{FALSE}, & \text{UNSAT}(P \wedge H) \\ \text{UNKNOWN}, & \text{otherwise.} \end{cases} \quad (1)$$

If premise consistency itself is violated (P is unsatisfiable) or any sentence fails to translate or compile, the item is flagged as a bridge failure rather than forced into one of the three labels.

D. Verification and Feedback Layer

A deterministic verifier (Algorithm 1) combines the neural prediction $\hat{y}_{\text{LLM}}$ and the symbolic verdict $\hat{y}_{\text{sym}}$. If the bridge failed to produce a verdict, the system ABSTAINs and reports the neural answer, explicitly unverified. If the two predictions agree, the system RETAINs the neural answer. If they disagree, the system CORRECTs the neural answer with the symbolic verdict. This design deliberately does not simply replace the LLM with the solver: when the bridge fails, the frozen LLM's answer is still returned, rather than the pipeline failing outright.

IV. EXPERIMENTAL METHODOLOGY

A. Research Questions

RQ1. Can an independent symbolic verification branch improve a frozen LLM's predictions on raw natural-language logical reasoning, without modifying the LLM?
RQ2. How reliable is the neural-to-symbolic bridge under controlled, synthetic language versus more complex, natural language?
RQ3. When the neural and symbolic branches disagree, does the resulting correction help or harm accuracy?

B. Datasets

We use **ProofWriter** [5], whose items consist of short, synthetic, closed-vocabulary facts and universally-quantified rules (e.g., "If something is cold then it sees the lion.") with minimal coreference, and **FOLIO** [6], whose items are natural, crowd-sourced sentences with richer syntax, cross-sentence entity references, and more varied quantifier and connective structure. Both are used exactly as originally distributed as raw text; no gold FOL or gold labels are ever supplied to either branch. FOLIO evaluation uses items from the project's frozen 30-item FOLIO prototype sample; ProofWriter evaluation uses a non-frozen development pilot for gating and a separate frozen 20-item pilot sample for the headline comparison (Section V), consistent with the project's own `proofwriter_phase1_samples=20` configuration.

C. Models

Branch A uses **Qwen3-8B** [2], loaded once and reused unmodified across all items, with greedy (temperature-0) decoding constrained to the label set {TRUE, FALSE, UNKNOWN}. Branch B uses the official **NL2Logic** implementation [1] with **Gemma-2-9B-Instruct** [3], quantized to `Q4_K_M` (4-bit) and served through **Ollama** [14], as its semantic-parsing model. Solving uses **Z3** [4] via the `z3-solver` [15] Python bindings, with a 10-second per-check timeout.

D. Hardware

All experiments ran on a single RTX 5050 Laptop GPU with approximately 8 GB of VRAM. The quantized Gemma-2-9B-Instruct model occupied approximately 6 GB and remained fully GPU-resident throughout; no CPU or shared-RAM offload was used or required.

E. Metrics

We report accuracy, macro-F1, per-class F1, AST validity rate, Z3-executability rate, verifier action counts (RETAIN/CORRECT/ABSTAIN), helpful and harmful correction counts, and per-sentence translation latency (median, mean, and maximum, given the sensitivity of the mean to outliers observed in our runs).

F. Leakage Prevention

Neither branch is ever given a gold label or gold logical form. Branch A's prompt contains only the raw facts, rules, and question/hypothesis. Branch B's translator receives only raw sentences; gold labels are used exclusively, after the fact, to score the already-produced predictions. The bridge itself (prompts, parser code, code generator) is the unmodified upstream NL2Logic implementation and was not tuned on any evaluation item.

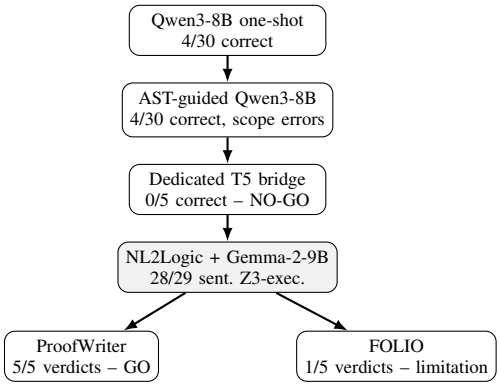


Fig. 2. Experimental progression of NL-to-FOL translation bridges evaluated in this study, and the controlled-language (ProofWriter) vs. natural-language (FOLIO) split observed once syntactic executability stopped being the bottleneck.

TABLE I
BRIDGE EVOLUTION ON FOLIO (SAMPLE SIZES DIFFER BY ROW)

Method (n)	Schema/AST valid	Z3-exec.	Verdict correct
Gold-FOL oracle (30)	—	—	29/30 (96.7%)
Qwen3-8B V1, one-shot (30)	9/30	8/30	4/30 (13.3%)
Qwen3-8B V2, AST-guided (30)	8/30	8/30	4/30 (13.3%)
T5 NL-to-FOL bridge (5)	3/5	3/5	0/5 (0%)
NL2Logic + Gemma-2-9B (5)	5/5 (29/29 sent.)	28/29 sent.	1/5 (20%)

TABLE II
FAILURE-MODE SHIFT BETWEEN TRANSLATION-V1 AND V2 ON THE FROZEN 30-ITEM FOLIO SAMPLE (DOMINANT ERROR CATEGORIES, SENTENCE-LEVEL COUNTS)

Error category	V1	V2
Malformed / schema error	32	2
Quantifier-scope error	—	32

TABLE III
FROZEN 20-ITEM PROOFWRITER PILOT: QWEN3-8B ALONE VS. FULL NSCE

Method	Accuracy	Macro-F1
Qwen3-8B only	40.0% (8/20)	0.300
Full NSCE	**70.0% (14/20)**	**0.676**

V. RESULTS

A. Bridge Evolution on FOLIO

Table I summarizes every FOL-translation approach evaluated over the course of this study on FOLIO, prior to and including the NL2Logic bridge. Sample sizes differ across rows and are reported explicitly rather than merged into one column, because the 30-item and 5-item samples are not directly comparable in statistical power.

The gold-FOL oracle row uses the officially provided logical forms, not a model-produced translation, and establishes that Z3 itself is reliable (96.7% agreement) when the formalization is correct. Translation-v1 (one-shot Qwen3-8B) and Translation-v2 (a decomposed, AST-inspired Qwen3-8B pipeline) both plateau at 4/30 correct; comparing their error categories (Table II) shows that v2 essentially traded one failure mode for another rather than improving semantic correctness. The dedicated T5 bridge (`fvossel/t5-base-nl-to-fol`) was stopped under a predefined 5-item NO-GO gate at 0/5 semantically correct. NL2Logic with Gemma-2-9B achieves near-complete syntactic executability (28/29 individual premise/hypothesis sentences, 96.6%) but only 1/5 correct final entailment verdicts – a large jump in syntactic quality that did not carry through to semantic correctness (Section VI).

B. FOLIO NL2Logic Gate

On the 5-item FOLIO gate, drawn from the frozen 30-item prototype and covering binary-relation/multi-variable, universal-implication-with-negation, existential, conjunction/disjunction, and negation-with-binary-relation constructs, translation latency had a median of 12.1 s and a maximum of 30.2 s per sentence. No sentence timed out or crashed; the low verdict accuracy is a semantic, not an operational, failure (Section VI).

C. Raw-Text ProofWriter Gate

On a 5-item, non-frozen ProofWriter development-pilot gate, the same unmodified bridge (only the item loader differs, mapping ProofWriter's facts/rules/question format instead of FOLIO's premises/hypothesis format) achieved 5/5 valid ASTs, 5/5 Z3-executable theories, and 5/5 correct verdicts, at a median latency of 11.6 s/sentence. This gate passing decisively is what licensed proceeding to the frozen 20-item pilot below, under the same predefined go/no-go protocol that stopped the T5 bridge and (separately) the FOLIO branch of this study.

D. Frozen 20-Item ProofWriter Pilot

Table III reports the headline comparison on the frozen 20-item ProofWriter pilot. We term this a *pilot*, not a benchmark evaluation, given $n = 20$.

Full NSCE improves accuracy by an absolute **+30 percentage points** over Qwen3-8B alone. On this sample, the bridge translated and Z3-solved all 20 items successfully (0 bridge failures), and the standalone symbolic-branch accuracy is also 14/20 (70%) – a relationship we return to in Section VI.

Table IV shows that of the 14 retained items, 8 were correct and 6 were cases where both branches agreed and were both wrong (an upper bound the verifier cannot fix, since it only acts on disagreement). All 6 disagreements were resolved by CORRECTing toward the symbolic verdict, and every one of the 6 corrections was helpful; none were harmful. Treating this as paired data (Qwen wrong $\rightarrow$ NSCE correct: 6; Qwen correct $\rightarrow$ NSCE wrong: 0), an exact McNemar (sign) test on the 6 discordant pairs gives a two-sided $p = 0.03125$ ($p = 2^{-5}$), which we report descriptively rather than as strong confirmatory evidence, given the small discordant-pair count. Wilson 95% confidence intervals on the two

TABLE IV
VERIFICATION ACTIONS ON THE FROZEN 20-ITEM PILOT

Action	Count
RETAIN	14
CORRECT	6
ABSTAIN	0
Helpful corrections	6
Harmful corrections	0

accuracy point estimates are wide and overlapping (Qwen-only: $[21.9\%, 61.3\%]$; full NSCE: $[48.1\%, 85.5\%]$), which we discuss as a direct consequence of $n = 20$ in Section VII.

E. Latency

On the frozen 20-item ProofWriter pilot, per-sentence translation latency had a median of 11.4 s and a mean of 22.9 s; the mean is pulled upward by a transient slowdown affecting four consecutive sentences during one stretch of the run, including one outlier of approximately 2019 s, coincident with a stall in the process's own progress monitor rather than with any particular sentence's content. We report the median as the representative figure and discuss the outlier as an operational limitation in Section VII, rather than treating the mean as representative.

VI. ANALYSIS AND DISCUSSION

Why ProofWriter works. ProofWriter's sentences are short, use a small closed vocabulary, and express one relation per sentence with explicit universal quantification ("If something is X then it Y"). Each fact or rule maps almost directly onto NL2Logic's per-sentence unary/binary predicate scheme, so entities and predicates referenced in different sentences of the same item are far more likely to be named consistently, letting Z3 reason over a theory that actually reflects the source text.

Why FOLIO does not. FOLIO's sentences are longer, more syntactically varied, and more often depend on how an earlier sentence introduced or described an entity. In our worked FOLIO gate example (idx 180, gold label FALSE), "Susan flies *to* LGA airport," "The departure and arrival can not be the same airport," and "John flies *from* LGA airport" each translate individually to syntactically valid, Z3-executable formulas, but the second sentence's constraint is captured as an unrelated unary predicate rather than as a relation linking the first and third sentences' airport arguments. The resulting theory is internally consistent but does not encode the cross-sentence constraint the premises actually assert, so Z3 correctly reports UNKNOWN on a theory that was never given enough information to entail the true answer. **Executable logic does not imply semantically faithful logic**: our syntactic-executability rate rose from roughly 27% (8/30) for both Qwen3-8B translators to 96.6% (28/29 sentences) for NL2Logic, while verdict correctness on the FOLIO gate did not rise correspondingly (4/30 $\rightarrow$ 1/5). The bottleneck moved from "can the LLM produce a parseable formula" to "does the produced formula preserve the source text's cross-sentence relations," and NL2Logic's per-sentence, per-clause parsing design – while effective at the former – does not by itself solve the latter, a limitation NL2Logic's own authors note under coreference handling [1].

The verifier tracks the symbolic branch when the bridge does not fail. On the frozen 20-item ProofWriter pilot, the bridge never failed (0 abstentions), so by construction every disagreement is resolved toward the symbolic verdict, and the standalone symbolic-branch accuracy (14/20) exactly equals the full system's accuracy (14/20). This is not a coincidence of the verifier policy being especially clever; it is a direct consequence of the bridge's reliability on this sample. Attributing the +30-point gain to the verification *policy* alone would overstate the contribution of the RETAIN/CORRECT/ABSTAIN logic itself and understate the dependency on bridge quality – on FOLIO, where the bridge is far less reliable, we would not expect the same policy to produce a comparable gain, and we did not test the full NSCE loop on FOLIO for this reason (Section VII).

VII. THREATS TO VALIDITY AND LIMITATIONS

- **Pilot sample size.** The headline comparison uses $n = 20$ frozen items. The resulting Wilson confidence intervals on accuracy are wide and overlapping, and the exact McNemar test is computed on only 6 discordant pairs; none of this should be read as a well-powered benchmark result.
- **ProofWriter is synthetic and closed-vocabulary.** The headline gain is demonstrated on language explicitly designed to be logically regular; we make no claim that it generalizes to less controlled text without further bridge work.
- **FOLIO gate is small (5 items).** The 1/5 verdict-correct result, while consistent with the underlying translation failure mode we identify with a worked example, is itself too small a sample to bound a FOLIO-wide accuracy figure; we therefore report it as a documented limitation rather than as a benchmark score for NL2Logic on FOLIO, which the original NL2Logic publication evaluates at larger scale and reports considerably higher performance on [1].
- **Model- and quantization-specific results.** All symbolic-branch results use one specific quantized checkpoint (Gemma-2-9B-Instruct, Q4_K_M) under one serving stack (Ollama) on one 8 GB GPU. We did not test unquantized Gemma-2-9B, larger Gemma variants, or alternative parsing models within NL2Logic.
- **Bridge dependence.** Because the bridge did not fail on the ProofWriter pilot, the full system's accuracy equals the symbolic branch's standalone accuracy on that sample; the RETAIN/CORRECT/ABSTAIN verifier's distinct contribution is better characterized by its zero-harmful-correction property than by the headline accuracy delta.
- **Latency and operational robustness.** A single-worker request queue combined with a purely advisory client-

side timeout meant that one transient slowdown (up to $\sim$2019 s for one sentence) did not abort the run but did inflate mean latency; a production system would need a hard, cancelling timeout and a multi-worker request pool.

- **No fine-tuning or larger-model comparison.** We compare a frozen 8B-parameter LLM against the same model augmented with verification; we do not compare against a fine-tuned Qwen3-8B or a larger frozen LLM (e.g., 70B-class), either of which might narrow or widen the observed gap.
- **Possible reproduction differences from the original NL2Logic FOLIO setup.** We use our own item selection, prompt harness, and Z3 aggregation logic around the unmodified NL2Logic library; differences from the original paper's exact FOLIO evaluation protocol (e.g., dataset preprocessing, aggregation across premises) may partly explain the gap between our 1/5 gate result and NL2Logic's own reported FOLIO numbers.

VIII. Future Work

Natural directions include: evaluating on a larger held-out ProofWriter and FOLIO split once bridge reliability on the latter improves; reproducing NL2Logic's exact published FOLIO configuration to isolate our harness's contribution to the gap; adding explicit shared entity/relation normalization across sentences of the same item (e.g., coreference resolution before parsing) to address the cross-sentence linking failure identified in Section VI; testing stronger or differently-quantized semantic-parsing models inside NL2Logic; calibrating the verifier's confidence rather than treating agreement/disagreement as binary; and extending bridge-confidence-aware abstention so the system can decline to correct when the symbolic branch itself shows signs of a degenerate translation (e.g., disconnected predicate graphs), rather than only abstaining on outright translation failure.

IX. Conclusion

On a frozen, small ProofWriter pilot, wiring an independent neural-to-symbolic-plus-solver verification branch around a frozen Qwen3-8B substantially improved its accuracy (40.0% $\rightarrow$ 70.0%) with zero observed harmful corrections, without any fine-tuning of the base model. This supports the premise that independent symbolic verification can meaningfully improve a frozen LLM's logical reliability, *when the translation bridge is reliable*. Our parallel FOLIO study shows that reliability is not guaranteed by the bridge's syntactic executability: NL2Logic with Gemma-2-9B produced valid, solver-executable logic on 96.6% of FOLIO sentences yet only 1 of 5 correct downstream verdicts, because cross-sentence semantic faithfulness – not solver correctness or syntactic validity – is where the system currently breaks on more natural language. We report both results together, as we believe an honest account of a verification architecture's limits is as important as its headline gain.

Acknowledgment

The authors thank the maintainers of NL2Logic, Z3, Qwen, Gemma, and Ollama for releasing the open models and tools this study depends on.

References

[1] R. R. Putra, R. S. P. Basuki, Y. Cheng, and P. Gao, "NL2LOGIC: AST-guided translation of natural language into first-order logic with large language models," *arXiv preprint arXiv:2602.13237*, 2026.

[2] Qwen Team, "Qwen3 technical report," *arXiv preprint arXiv:2505.09388*, 2025.

[3] Gemma Team, "Gemma 2: Improving open language models at a practical size," *arXiv preprint arXiv:2408.00118*, 2024.

[4] L. De Moura and N. Bjørner, "Z3: An efficient SMT solver," in *Proceedings of the Theory and Practice of Software, 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'08/ETAPS'08)*. Berlin, Heidelberg: Springer-Verlag, 2008, pp. 337–340.

[5] O. Tafjord, B. Dalvi, and P. Clark, "ProofWriter: Generating implications, proofs, and abductive statements over natural language," in *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*. Online: Association for Computational Linguistics, 2021, pp. 3621–3634.

[6] S. Han, H. Schoelkopf, Y. Zhao, Z. Qi, M. Riddell, W. Zhou, J. Coady, D. Peng, Y. Qiao, L. Benson *et al.*, "FOLIO: Natural language reasoning with first-order logic," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024, pp. 22 017–22 031.

[7] M. K. Sarker, L. Zhou, A. Eberhart, and P. Hitzler, "Towards data- and knowledge-driven artificial intelligence: A survey on neuro-symbolic computing," *arXiv preprint arXiv:2210.15889*, 2022.

[8] L. Pan, A. Albalak, X. Wang, and W. Wang, "Logic-LM: Empowering large language models with symbolic solvers for faithful logical reasoning," in *Findings of the Association for Computational Linguistics: EMNLP 2023*. Singapore: Association for Computational Linguistics, 2023, pp. 3806–3824.

[9] T. Olausson, A. Gu, B. Lipkin, C. Zhang, A. Solar-Lezama, J. Tenenbaum, and R. Levy, "LINC: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Singapore: Association for Computational Linguistics, 2023, pp. 5153–5176.

[10] F. Raspanti, T. Ozcelebi, and M. Holenderski, "Grammar-constrained decoding makes large language models better logical parsers," in *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 6: Industry Track)*, 2025, pp. 485–499.

[11] J. Liu, "Few-shot natural language to first-order logic translation via code generation," in *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Albuquerque, New Mexico: Association for Computational Linguistics, 2025, pp. 10 939–10 960.

[12] J. Tian, Y. Li, W. Chen, L. Xiao, H. He, and Y. Jin, "Diagnosing the first-order logical reasoning ability through LogicNLI," in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 2021, pp. 3738–3747.

[13] J. Wei, X. Wang, D. Schuurmans, M. Bosma, F. Xia, E. Chi, Q. V. Le, D. Zhou *et al.*, "Chain-of-thought prompting elicits reasoning in large language models," *Advances in Neural Information Processing Systems*, vol. 35, pp. 24 824–24 837, 2022.

[14] Ollama Inc., "Ollama: Get up and running with large language models locally," https://ollama.com, 2026.

[15] Z3 Prover Contributors, "z3-solver: Python bindings for the Z3 theorem prover," PyPI, 2026, https://pypi.org/project/z3-solver/.